

Raashid Salih -MLOps Engineer-

✉ rashidsaleh11@gmail.com 📞 +91 62-826-22077 🔗 raashidsalih.github.io

PROFESSIONAL EXPERIENCE

MLOps Engineer

Jun 2025 – Present | Singapore

CoolRIOTS (Business Execution Platform)

- **Led the deployment and management of on-premises LLMs**, including developing an asynchronous engine implementation for inference, streaming, and batching, and driving full-stack integration.
- **Crafted a robust benchmarking suite** to fine-tune vLLM parameters for ideal resource utilization.
- **Engineered a high-speed model registry** utilizing Redis for in-memory caching of model endpoint details.
- **Architected a multi-cluster ArgoCD environment** to ensure continuous deployment of production services.
- **Developed and deployed robust S3-based loading solutions** for large models, ensuring reliable and efficient data access across the platform.
- **Involved with Hashicorp Vault platform architecture** to reduce system latency while maintaining security.
- **Established a comprehensive knowledge base** for LLM/vLLM pertinent configurations to accelerate deployment cycles.

Data Scientist

Dec 2023 – Jun 2025 | Bangalore, India

FreightFox (Logistics Startup) 🔗

- Shaped company's **Data Science strategy** by aligning development with product roadmap for maximum impact.
- **Optimized production ready Python and SQL code** to be able to handle large datasets under limited hardware.
- Engineered **industry-first latency index** to highlight network latency changes and aid clients with planning.
- Created **comprehensive anomaly detection suite** which reduces client cognitive burden by highlighting critical actionable shipments.
- Developed **RAG based LLM solution (Text2SQL)** to facilitate self serve analytics via chat interface.
- **Markedly enhanced data quality** by implementing novel solutions to industry-wide problems of non-standard location and vehicle names, enabling inter-manufacturer analytics.
- **Productionized demand forecasting pipeline** with monitoring and evaluation to aid client in inventory management.
- **Developed an optimization model to minimize costs** by determining an optimal fleet mix for any given payload.

Data & Analytics Intern

Jan 2022 – Sep 2022 | KL, Malaysia

Schlumberger (Energy MNC) 🔗

- **Enhanced predictive model reliability** by applying **Explainable ML techniques** leading to identification of model strengths and weaknesses.
- **Implemented robust pipelines for sprint planning and department KPO dashboards** ensuring comprehensive documentation and adherence to best practices, improving data visibility.

SKILLS

Development & Databases

Python, FastAPI, Pydantic, PostgreSQL, Redis, dbt, Postman, PlayWright, Ruff, Bash

DevOps & Orchestration

AWS, Openshift, Docker, Kubernetes, ArgoCD, Airflow, Terraform, MLFlow, Grafana, Helm, Vault

GenAI and Machine Learning

vLLM, OpenVINO, KServe, Qdrant, Streamlit, Superset, Dash, PuLP, Metabase, AutoML

PROJECTS

End-to-End Churn Analytics Platform 🔗

A cloud-hosted containerized data platform to simulate and analyze customer behavior utilizing microservices architecture. This platform leverages MLOps best practices, including MLFlow and Grafana for experiment tracking and monitoring, and implements CI/CD for seamless updates.

Agmarknet Scraper 🔗

A CLI tool designed to gather data from Agmarknet 🔗, employing a direct URL substitution approach for stability and built with an MVC architectural pattern for better separation of concerns.

JenksPy: Fisher-Jenks algorithm (Python) 🔗

Open source contribution to help determine ideal cluster size using elbow charts based on the goodness of variance fit (GVF) metric.

EDUCATION

Georgia Institute of Technology (GaTech) [On Academic Hiatus]

2023 – present | Atlanta, Georgia, USA

Master of Science in Analytics (MSc)

- CGPA - 4.00/4.00

International Islamic University Malaysia (IIUM)

2018 – 2022 | Kuala Lumpur, Malaysia

Bachelor of Computer Science (Honours)

- CGPA - 3.99/4.00

IELTS Academic: Band 9 | CEFR - C2 (highest level of English proficiency)